

Distributed Systems

The second half of *Concurrent and Distributed Systems*

<https://www.cl.cam.ac.uk/teaching/current/ConcDisSys>

Dr. Martin Kleppmann (mk428@cam)

University of Cambridge

Computer Science Tripos, Part IB



This work is published under a
Creative Commons BY-SA license.

Lecture 5

Replication

Replication

- ▶ Keeping a copy of the same data on multiple nodes
- ▶ Databases, filesystems, caches, ...
- ▶ A node that has a copy of the data is called a **replica**

Replication

- ▶ Keeping a copy of the same data on multiple nodes
- ▶ Databases, filesystems, caches, ...
- ▶ A node that has a copy of the data is called a **replica**
- ▶ If some replicas are faulty, others are still accessible
- ▶ Spread load across many replicas

Replication

- ▶ Keeping a copy of the same data on multiple nodes
- ▶ Databases, filesystems, caches, ...
- ▶ A node that has a copy of the data is called a **replica**
- ▶ If some replicas are faulty, others are still accessible
- ▶ Spread load across many replicas
- ▶ Easy if the data doesn't change: just copy it
- ▶ We will focus on data changes

Replication

- ▶ Keeping a copy of the same data on multiple nodes
- ▶ Databases, filesystems, caches, ...
- ▶ A node that has a copy of the data is called a **replica**
- ▶ If some replicas are faulty, others are still accessible
- ▶ Spread load across many replicas
- ▶ Easy if the data doesn't change: just copy it
- ▶ We will focus on data changes

Compare to **RAID** (Redundant Array of Independent Disks):
replication within a single computer

- ▶ RAID has single controller; in distributed system, each node acts independently
- ▶ Replicas can be distributed around the world, near users


Retrying state updates

User A: The moon is not actually made of cheese!




12,300 people like this.

client




Retrying state updates

User A: The moon is not actually made of cheese!





12,300 people like this.



Retrying state updates

User A: The moon is not actually made of cheese!

 Like 12,300 people like this.




Retrying state updates

User A: The moon is not actually made of cheese!



12,300 people like this.



Deduplicating requests requires that the database tracks which requests it has already seen (in stable storage)



TWEETS
6,219

FOLLOWING
-20

FOLLOWERS
24.1K



 **Follow**

Лепра

@lepratorium



Добро пожаловать отсюда

Default City



Лепра @lepratorium · 2h

Викторианские советы

Часть 2 pic.twitter.com/21PraRYBaO

Details



Лепра @lepratorium · 2h

Викторианские советы

Часть 1 pic.twitter.com/BVE6ao8711

Details

[Go to full profile](#)

Idempotence

A function f is idempotent if $f(x) = f(f(x))$.

- ▶ **Not idempotent:** $f(\text{likeCount}) = \text{likeCount} + 1$
- ▶ **Idempotent:** $f(\text{likeSet}) = \text{likeSet} \cup \{\text{userID}\}$

Idempotent requests can be retried without deduplication.

Idempotence

A function f is idempotent if $f(x) = f(f(x))$.

- ▶ **Not idempotent:** $f(\text{likeCount}) = \text{likeCount} + 1$
- ▶ **Idempotent:** $f(\text{likeSet}) = \text{likeSet} \cup \{\text{userID}\}$


Idempotent requests can be retried without deduplication.

Choice of retry behaviour:

- ▶ **At-most-once** semantics:
send request, don't retry, update may not happen
- ▶ **At-least-once** semantics:
retry request until acknowledged, may repeat update
- ▶ **Exactly-once** semantics:
retry + idempotence or deduplication


Adding and then removing again

client 1



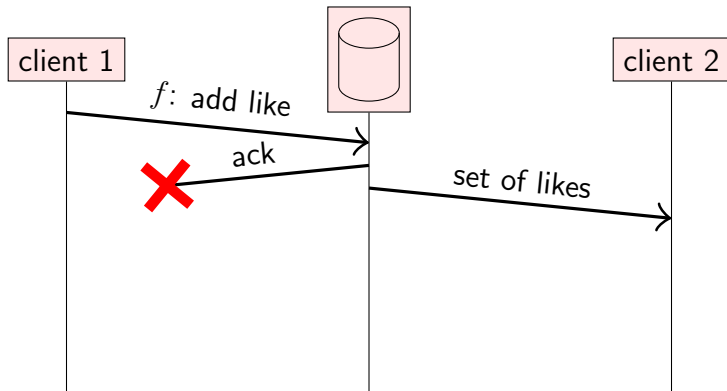
client 2

Adding and then removing again



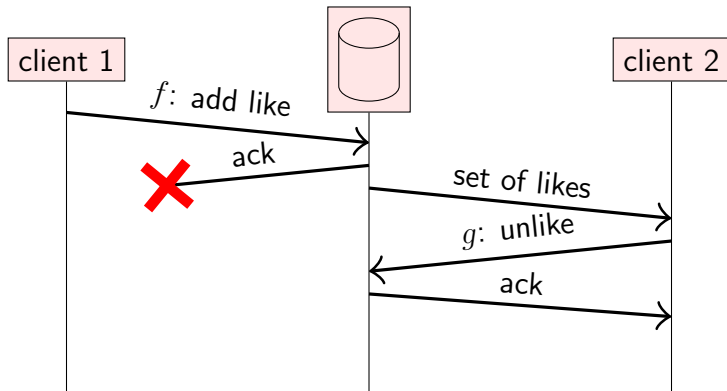
$$f(\text{likes}) = \text{likes} \cup \{\text{userID}\}$$

Adding and then removing again



$$f(\text{likes}) = \text{likes} \cup \{\text{userID}\}$$


Adding and then removing again



$$f(\text{likes}) = \text{likes} \cup \{\text{userID}\}$$

$$g(\text{likes}) = \text{likes} \setminus \{\text{userID}\}$$


Adding and then removing again



$$f(\text{likes}) = \text{likes} \cup \{\text{userID}\}$$

$$g(\text{likes}) = \text{likes} \setminus \{\text{userID}\}$$

Adding and then removing again




$$f(\text{likes}) = \text{likes} \cup \{\text{userID}\}$$


$$g(\text{likes}) = \text{likes} \setminus \{\text{userID}\}$$

Idempotent? $f(f(x)) = f(x)$ but $f(g(f(x))) \neq g(f(x))$


Another problem with adding and removing




Another problem with adding and removing



Another problem with adding and removing




Another problem with adding and removing




Final state ($x \notin A$, $x \in B$) is the same as in this case:


Another problem with adding and removing



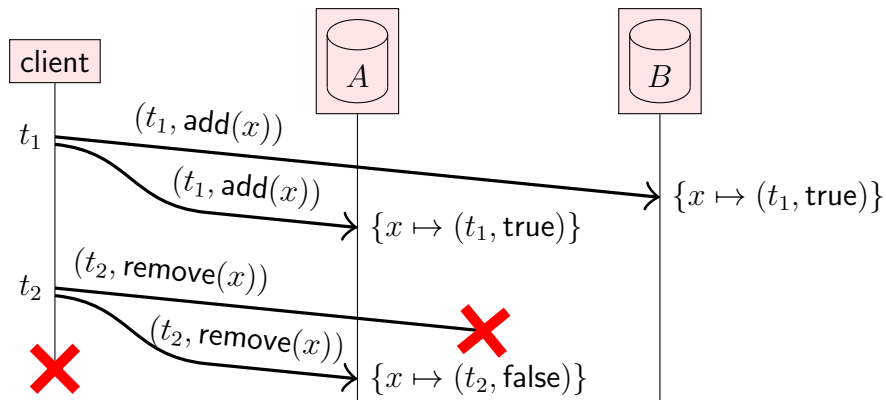
Final state ($x \notin A$, $x \in B$) is the same as in this case:




Timestamps and tombstones



Timestamps and tombstones

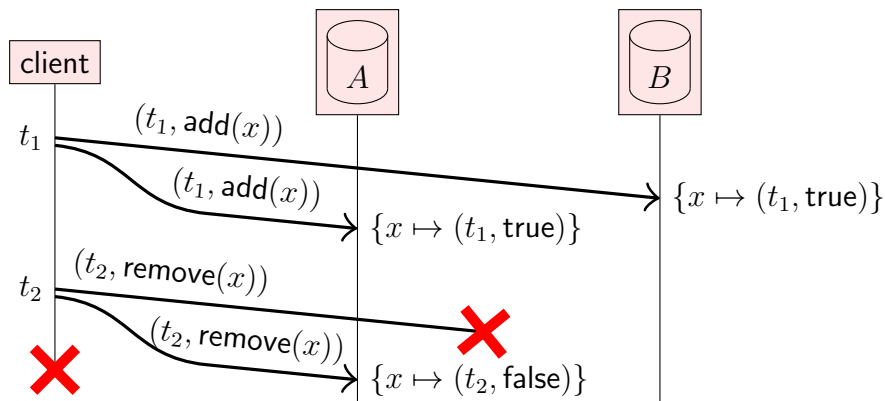


Timestamps and tombstones



“remove(x)” doesn’t actually remove x : it labels x with “false” to indicate it is invisible (a **tombstone**)

Timestamps and tombstones





“remove(x)” doesn’t actually remove x : it labels x with “false” to indicate it is invisible (a **tombstone**)

Every record has **logical timestamp** of last write




Reconciling replicas

Replicas periodically communicate among themselves to check for any inconsistencies.


$$\{x \mapsto (t_2, \text{false})\}$$

$$\{x \mapsto (t_1, \text{true})\}$$




Reconciling replicas

Replicas periodically communicate among themselves to check for any inconsistencies.




Reconciling replicas

Replicas periodically communicate among themselves to check for any inconsistencies.




Reconciling replicas

Replicas periodically communicate among themselves to check for any inconsistencies.




Propagate the record with the latest timestamp,
discard the records with earlier timestamps
(for a given key).


Concurrent writes by different clients




Concurrent writes by different clients



Concurrent writes by different clients



Concurrent writes by different clients




Two common approaches:

- **Last writer wins (LWW):**

Use timestamps with total order (e.g. Lamport clock)

Keep v_2 and discard v_1 if $t_2 > t_1$. Note: **data loss!**

Concurrent writes by different clients



Two common approaches:

- ▶ **Last writer wins (LWW):**

Use timestamps with total order (e.g. Lamport clock)
Keep v_2 and discard v_1 if $t_2 > t_1$. Note: **data loss!**

- ▶ **Multi-value register:**

Use timestamps with partial order (e.g. vector clock)
 v_2 replaces v_1 if $t_2 > t_1$; preserve both $\{v_1, v_2\}$ if $t_1 \parallel t_2$

Probability of faults

A replica may be **unavailable** due to network partition or node fault (e.g. crash, hardware problem).

Probability of faults

A replica may be **unavailable** due to network partition or node fault (e.g. crash, hardware problem).

Assume each replica has probability p of being faulty or unavailable at any one time, and that faults are independent.
(Not actually true! But okay approximation for now.)

Probability of faults

A replica may be **unavailable** due to network partition or node fault (e.g. crash, hardware problem).

Assume each replica has probability p of being faulty or unavailable at any one time, and that faults are independent.
(Not actually true! But okay approximation for now.)

Probability of **all** n replicas being faulty: p^n

Probability of ≥ 1 out of n replicas being faulty: $1 - (1 - p)^n$

Probability of faults

A replica may be **unavailable** due to network partition or node fault (e.g. crash, hardware problem).

Assume each replica has probability p of being faulty or unavailable at any one time, and that faults are independent.
(Not actually true! But okay approximation for now.)


Probability of **all** n replicas being faulty: p^n

Probability of ≥ 1 out of n replicas being faulty: $1 - (1 - p)^n$


Example with $p = 0.01$:

replicas n	$P(\geq 1 \text{ faulty})$	$P(\geq \frac{n+1}{2} \text{ faulty})$	$P(\text{all } n \text{ faulty})$
1	0.01	0.01	0.01
3	0.03	$3 \cdot 10^{-4}$	10^{-6}
5	0.049	$1 \cdot 10^{-5}$	10^{-10}
100	0.63	$6 \cdot 10^{-74}$	10^{-200}


Read-after-write consistency



Read-after-write consistency




Read-after-write consistency



Writing to one replica, reading from another: client does not read back the value it has written


Read-after-write consistency




Writing to one replica, reading from another: client does not read back the value it has written

Require writing to/reading from both replicas \implies cannot write/read if one replica is unavailable


Quorum (2 out of 3)



Quorum (2 out of 3)




Quorum (2 out of 3)




Write succeeds on B and C

Quorum (2 out of 3)




Write succeeds on B and C

Quorum (2 out of 3)



Write succeeds on B and C ; read succeeds on A and B

Quorum (2 out of 3)



Write succeeds on B and C; read succeeds on A and B
Choose between (t_0, v_0) and (t_1, v_1) based on timestamp

Read and write quorums

In a system with n replicas:

- ▶ If a write is acknowledged by w replicas (**write quorum**),

Read and write quorums

In a system with n replicas:

- ▶ If a write is acknowledged by w replicas (**write quorum**),
- ▶ and we subsequently read from r replicas (**read quorum**),
- ▶ and $r + w > n$,

Read and write quorums


In a system with n replicas:

- ▶ If a write is acknowledged by w replicas (**write quorum**),
- ▶ and we subsequently read from r replicas (**read quorum**),
- ▶ and $r + w > n$,
- ▶ ...then the read will see the previously written value (or a value that subsequently overwrote it)

Read and write quorums

In a system with n replicas:


- ▶ If a write is acknowledged by w replicas (**write quorum**),
- ▶ and we subsequently read from r replicas (**read quorum**),
- ▶ and $r + w > n$,
- ▶ ... then the read will see the previously written value (or a value that subsequently overwrote it)
- ▶ Read quorum and write quorum share ≥ 1 replica




Read and write quorums

In a system with n replicas:


- ▶ If a write is acknowledged by w replicas (**write quorum**),
- ▶ and we subsequently read from r replicas (**read quorum**),
- ▶ and $r + w > n$,
- ▶ ... then the read will see the previously written value (or a value that subsequently overwrote it)
- ▶ Read quorum and write quorum share ≥ 1 replica
- ▶ Typical: $r = w = \frac{n+1}{2}$ for $n = 3, 5, 7, \dots$ (majority)
- ▶ Reads can tolerate $n - r$ unavailable replicas, writes $n - w$




Read repair



Read repair




Read repair



Update (t_1, v_1) is more recent than (t_0, v_0) since $t_0 < t_1$.

Read repair



Update (t_1, v_1) is more recent than (t_0, v_0) since $t_0 < t_1$.
Client helps **propagate** (t_1, v_1) to other replicas.

State machine replication

So far we have used best-effort broadcast for replication.
What about stronger broadcast models?

State machine replication

So far we have used best-effort broadcast for replication.
What about stronger broadcast models?

Total order broadcast: every node delivers the **same messages** in the **same order**

State machine replication

So far we have used best-effort broadcast for replication.
What about stronger broadcast models?

Total order broadcast: every node delivers the **same messages** in the **same order**

State machine replication (SMR):

- ▶ FIFO-total order broadcast every update to all replicas
- ▶ Replica delivers update message: apply it to own state

State machine replication

So far we have used best-effort broadcast for replication.
What about stronger broadcast models?

Total order broadcast: every node delivers the **same messages** in the **same order**

State machine replication (SMR):

- ▶ FIFO-total order broadcast every update to all replicas
- ▶ Replica delivers update message: apply it to own state
- ▶ Applying an update is deterministic

State machine replication

So far we have used best-effort broadcast for replication.
What about stronger broadcast models?

Total order broadcast: every node delivers the **same messages** in the **same order**

State machine replication (SMR):

- ▶ FIFO-total order broadcast every update to all replicas
- ▶ Replica delivers update message: apply it to own state
- ▶ Applying an update is deterministic
- ▶ Replica is a **state machine**: starts in fixed initial state, goes through same sequence of state transitions in the same order \implies all replicas end up in the same state

State machine replication

on request to perform update u **do**
 send u via FIFO-total order broadcast
end on

on delivering u through FIFO-total order broadcast **do**
 update state using arbitrary deterministic logic!
end on

State machine replication

on request to perform update u **do**
 send u via FIFO-total order broadcast
end on

on delivering u through FIFO-total order broadcast **do**
 update state using arbitrary deterministic logic!
end on

Closely related ideas:

- ▶ Serializable transactions (execute in delivery order)

State machine replication

on request to perform update u **do**
 send u via FIFO-total order broadcast
end on

on delivering u through FIFO-total order broadcast **do**
 update state using arbitrary deterministic logic!
end on

Closely related ideas:

- ▶ Serializable transactions (execute in delivery order)
- ▶ Blockchains, distributed ledgers, smart contracts

State machine replication

on request to perform update u **do**
 send u via FIFO-total order broadcast
end on

on delivering u through FIFO-total order broadcast **do**
 update state using arbitrary deterministic logic!
end on

Closely related ideas:

- ▶ Serializable transactions (execute in delivery order)
- ▶ Blockchains, distributed ledgers, smart contracts

Limitations:

- ▶ Cannot update state immediately, have to wait for delivery through broadcast

State machine replication

```
on request to perform update  $u$  do  
    send  $u$  via FIFO-total order broadcast  
end on
```

```
on delivering  $u$  through FIFO-total order broadcast do  
    update state using arbitrary deterministic logic!  
end on
```

Closely related ideas:


- ▶ Serializable transactions (execute in delivery order)
- ▶ Blockchains, distributed ledgers, smart contracts

Limitations:

- ▶ Cannot update state immediately, have to wait for delivery through broadcast
- ▶ Need fault-tolerant total order broadcast: see lecture 6


Database leader replica

Leader database replica L ensures total order broadcast




Database leader replica

Leader database replica L ensures total order broadcast




Database leader replica

Leader database replica L ensures total order broadcast



Database leader replica


Leader database replica L ensures total order broadcast



Follower F applies transaction log in commit order

Database leader replica

Leader database replica L ensures total order broadcast



Follower F applies transaction log in commit order

Replication using causal (and weaker) broadcast

State machine replication uses (FIFO-)total order broadcast.
Can we use weaker forms of broadcast too?

Replication using causal (and weaker) broadcast

State machine replication uses (FIFO-)total order broadcast.
Can we use weaker forms of broadcast too?

If replica state updates are **commutative**, replicas can process updates in different orders and still end up in the same state.

Updates f and g are commutative if $f(g(x)) = g(f(x))$

Replication using causal (and weaker) broadcast

State machine replication uses (FIFO-)total order broadcast.
Can we use weaker forms of broadcast too?

If replica state updates are **commutative**, replicas can process updates in different orders and still end up in the same state.

Updates f and g are commutative if $f(g(x)) = g(f(x))$

broadcast	assumptions about state update function
total order	deterministic (SMR)

Replication using causal (and weaker) broadcast

State machine replication uses (FIFO-)total order broadcast.
Can we use weaker forms of broadcast too?

If replica state updates are **commutative**, replicas can process updates in different orders and still end up in the same state.

Updates f and g are commutative if $f(g(x)) = g(f(x))$

broadcast	assumptions about state update function
total order	deterministic (SMR)
causal	deterministic, concurrent updates commute

Replication using causal (and weaker) broadcast

State machine replication uses (FIFO-)total order broadcast.
Can we use weaker forms of broadcast too?

If replica state updates are **commutative**, replicas can process updates in different orders and still end up in the same state.

Updates f and g are commutative if $f(g(x)) = g(f(x))$

broadcast	assumptions about state update function
total order	deterministic (SMR)
causal	deterministic, concurrent updates commute
reliable	deterministic, all updates commute

Replication using causal (and weaker) broadcast

State machine replication uses (FIFO-)total order broadcast.
Can we use weaker forms of broadcast too?

If replica state updates are **commutative**, replicas can process updates in different orders and still end up in the same state.

Updates f and g are commutative if $f(g(x)) = g(f(x))$

broadcast	assumptions about state update function
total order	deterministic (SMR)
causal	deterministic, concurrent updates commute
reliable	deterministic, all updates commute
best-effort	deterministic, commutative, idempotent, tolerates message loss