

Distributed Systems

The second half of *Concurrent and Distributed Systems*

<https://www.cl.cam.ac.uk/teaching/current/ConcDisSys>

Dr. Martin Kleppmann (mk428@cam)

University of Cambridge

Computer Science Tripos, Part IB



This work is published under a
Creative Commons BY-SA license.

Lecture 6

Consensus

Fault-tolerant total order broadcast

Total order broadcast is very useful for state machine replication.

Can implement total order broadcast by sending all messages via a single **leader**.

Problem: what if leader crashes/becomes unavailable?

Fault-tolerant total order broadcast

Total order broadcast is very useful for state machine replication.

Can implement total order broadcast by sending all messages via a single **leader**.

Problem: what if leader crashes/becomes unavailable?

- ▶ **Manual failover:** a human operator chooses a new leader, and reconfigures each node to use new leader

Used in many databases! Fine for planned maintenance.

Fault-tolerant total order broadcast

Total order broadcast is very useful for state machine replication.

Can implement total order broadcast by sending all messages via a single **leader**.

Problem: what if leader crashes/becomes unavailable?

- ▶ **Manual failover:** a human operator chooses a new leader, and reconfigures each node to use new leader

Used in many databases! Fine for planned maintenance.

Unplanned outage? Humans are slow, may take a long time until system recovers. . .

Fault-tolerant total order broadcast

Total order broadcast is very useful for state machine replication.

Can implement total order broadcast by sending all messages via a single **leader**.

Problem: what if leader crashes/becomes unavailable?

- ▶ **Manual failover**: a human operator chooses a new leader, and reconfigures each node to use new leader

Used in many databases! Fine for planned maintenance.

Unplanned outage? Humans are slow, may take a long time until system recovers. . .

- ▶ Can we **automatically choose a new leader**?

Consensus and total order broadcast

- ▶ Traditional formulation of consensus: several nodes want to come to **agreement** about a single **value**

Consensus and total order broadcast

- ▶ Traditional formulation of consensus: several nodes want to come to **agreement** about a single **value**
- ▶ In context of total order broadcast: this value is the **next message to deliver**

Consensus and total order broadcast

- ▶ Traditional formulation of consensus: several nodes want to come to **agreement** about a single **value**
- ▶ In context of total order broadcast: this value is the **next message to deliver**
- ▶ Once one node **decides** on a certain message order, all nodes will decide the same order

Consensus and total order broadcast

- ▶ Traditional formulation of consensus: several nodes want to come to **agreement** about a single **value**
- ▶ In context of total order broadcast: this value is the **next message to deliver**
- ▶ Once one node **decides** on a certain message order, all nodes will decide the same order
- ▶ Consensus and total order broadcast are formally equivalent

Consensus and total order broadcast

- ▶ Traditional formulation of consensus: several nodes want to come to **agreement** about a single **value**
- ▶ In context of total order broadcast: this value is the **next message to deliver**
- ▶ Once one node **decides** on a certain message order, all nodes will decide the same order
- ▶ Consensus and total order broadcast are formally equivalent

Common consensus algorithms:

- ▶ **Paxos**: single-value consensus
- ▶ **Multi-Paxos**: generalisation to total order broadcast

Consensus and total order broadcast

- ▶ Traditional formulation of consensus: several nodes want to come to **agreement** about a single **value**
- ▶ In context of total order broadcast: this value is the **next message to deliver**
- ▶ Once one node **decides** on a certain message order, all nodes will decide the same order
- ▶ Consensus and total order broadcast are formally equivalent

Common consensus algorithms:

- ▶ **Paxos**: single-value consensus
Multi-Paxos: generalisation to total order broadcast
- ▶ **Raft, Viewstamped Replication, Zab**:
FIFO-total order broadcast by default

Consensus system models

Paxos, Raft, etc. assume a **partially synchronous, crash-recovery** system model.

Consensus system models

Paxos, Raft, etc. assume a **partially synchronous, crash-recovery** system model.

Why not asynchronous?

- ▶ **FLP result** (Fischer, Lynch, Paterson):

There is no deterministic consensus algorithm that is guaranteed to terminate in an asynchronous crash-stop system model.

Consensus system models

Paxos, Raft, etc. assume a **partially synchronous, crash-recovery** system model.

Why not asynchronous?

- ▶ **FLP result** (Fischer, Lynch, Paterson):
There is no deterministic consensus algorithm that is guaranteed to terminate in an asynchronous crash-stop system model.
- ▶ Paxos, Raft, etc. use clocks only used for timeouts/failure detector to ensure progress. Safety (correctness) does not depend on timing.

Consensus system models

Paxos, Raft, etc. assume a **partially synchronous, crash-recovery** system model.

Why not asynchronous?

- ▶ **FLP result** (Fischer, Lynch, Paterson):
There is no deterministic consensus algorithm that is guaranteed to terminate in an asynchronous crash-stop system model.
- ▶ Paxos, Raft, etc. use clocks only used for timeouts/failure detector to ensure progress. Safety (correctness) does not depend on timing.

There are also consensus algorithms for a partially synchronous **Byzantine** system model (used in blockchains)

Leader election

Multi-Paxos, Raft, etc. use a leader to sequence messages.

Leader election

Multi-Paxos, Raft, etc. use a leader to sequence messages.

- ▶ Use a **failure detector** (timeout) to determine suspected crash or unavailability of leader.
- ▶ On suspected leader crash, **elect a new one**.

Leader election

Multi-Paxos, Raft, etc. use a leader to sequence messages.

- ▶ Use a **failure detector** (timeout) to determine suspected crash or unavailability of leader.
- ▶ On suspected leader crash, **elect a new one**.
- ▶ Prevent **two leaders at the same time** (“split-brain”)!

Leader election

Multi-Paxos, Raft, etc. use a leader to sequence messages.

- ▶ Use a **failure detector** (timeout) to determine suspected crash or unavailability of leader.
- ▶ On suspected leader crash, **elect a new one**.
- ▶ Prevent **two leaders at the same time** (“split-brain”)!

Ensure ≤ 1 leader per **term**:

- ▶ Term is incremented every time a leader election is started


Leader election

Multi-Paxos, Raft, etc. use a leader to sequence messages.

- ▶ Use a **failure detector** (timeout) to determine suspected crash or unavailability of leader.
- ▶ On suspected leader crash, **elect a new one**.
- ▶ Prevent **two leaders at the same time** (“split-brain”)!

Ensure ≤ 1 leader per **term**:

- ▶ Term is incremented every time a leader election is started
- ▶ A node can only **vote once** per term
- ▶ Require a **quorum** of nodes to elect a leader in a term




Leader election

Multi-Paxos, Raft, etc. use a leader to sequence messages.

- ▶ Use a **failure detector** (timeout) to determine suspected crash or unavailability of leader.
- ▶ On suspected leader crash, **elect a new one**.
- ▶ Prevent **two leaders at the same time** (“split-brain”)!

Ensure ≤ 1 leader per **term**:

- ▶ Term is incremented every time a leader election is started
- ▶ A node can only **vote once** per term
- ▶ Require a **quorum** of nodes to elect a leader in a term




Leader election

Multi-Paxos, Raft, etc. use a leader to sequence messages.

- ▶ Use a **failure detector** (timeout) to determine suspected crash or unavailability of leader.
- ▶ On suspected leader crash, **elect a new one**.
- ▶ Prevent **two leaders at the same time** (“split-brain”)!

Ensure ≤ 1 leader per **term**:

- ▶ Term is incremented every time a leader election is started
- ▶ A node can only **vote once** per term
- ▶ Require a **quorum** of nodes to elect a leader in a term



Can we guarantee there is only one leader?

Can guarantee unique leader **per term**.

Can we guarantee there is only one leader?

Can guarantee unique leader **per term**.


Cannot prevent having multiple leaders from different terms.

Can we guarantee there is only one leader?

Can guarantee unique leader **per term**.

Cannot prevent having multiple leaders from different terms.

Example: node 1 is leader in term t , but due to a network partition it can no longer communicate with nodes 2 and 3:




Nodes 2 and 3 may elect a new leader in term $t + 1$.

Can we guarantee there is only one leader?

Can guarantee unique leader **per term**.

Cannot prevent having multiple leaders from different terms.

Example: node 1 is leader in term t , but due to a network partition it can no longer communicate with nodes 2 and 3:




Nodes 2 and 3 may elect a new leader in term $t + 1$.

Node 1 may not even know that a new leader has been elected!

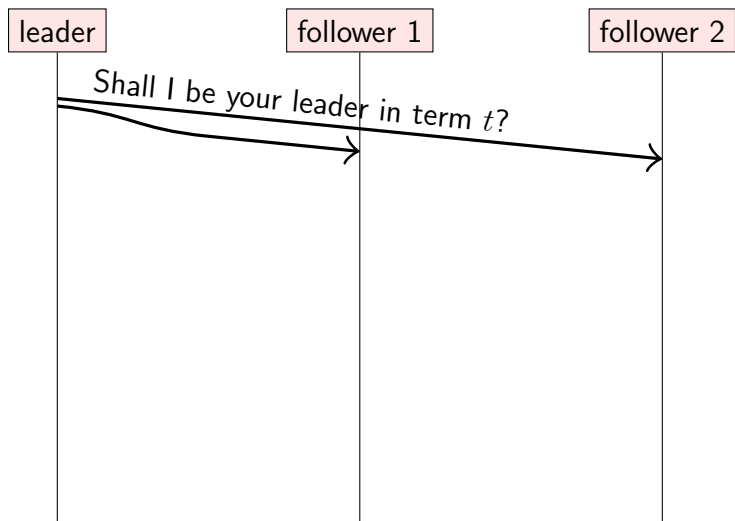
Checking if a leader has been voted out

For every decision (message to deliver), the leader must first get acknowledgements from a quorum.



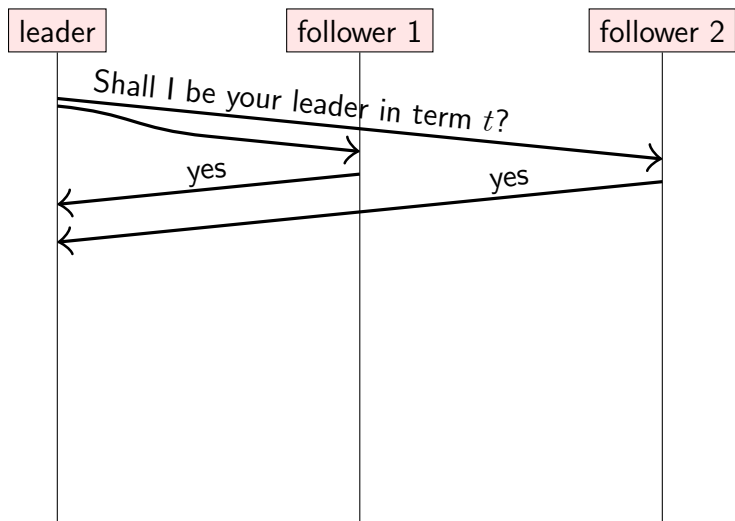
Checking if a leader has been voted out

For every decision (message to deliver), the leader must first get acknowledgements from a quorum.



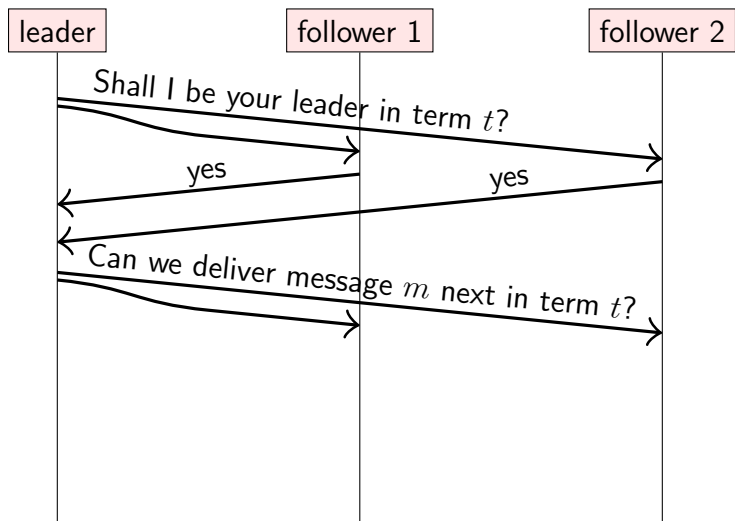
Checking if a leader has been voted out

For every decision (message to deliver), the leader must first get acknowledgements from a quorum.



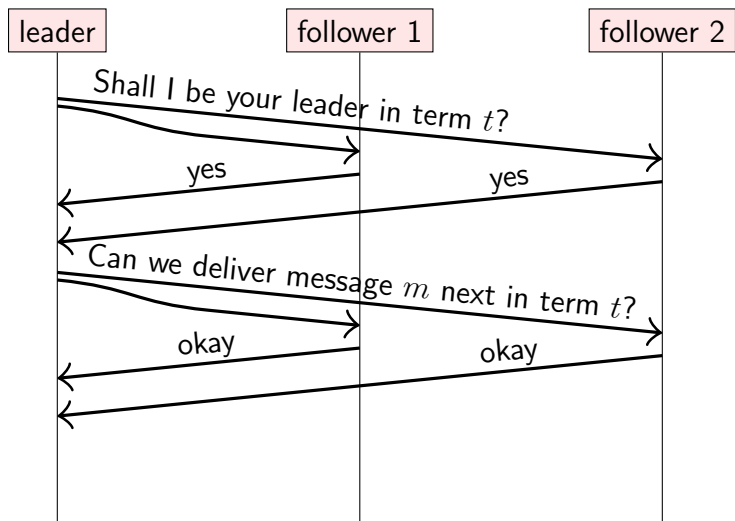
Checking if a leader has been voted out

For every decision (message to deliver), the leader must first get acknowledgements from a quorum.



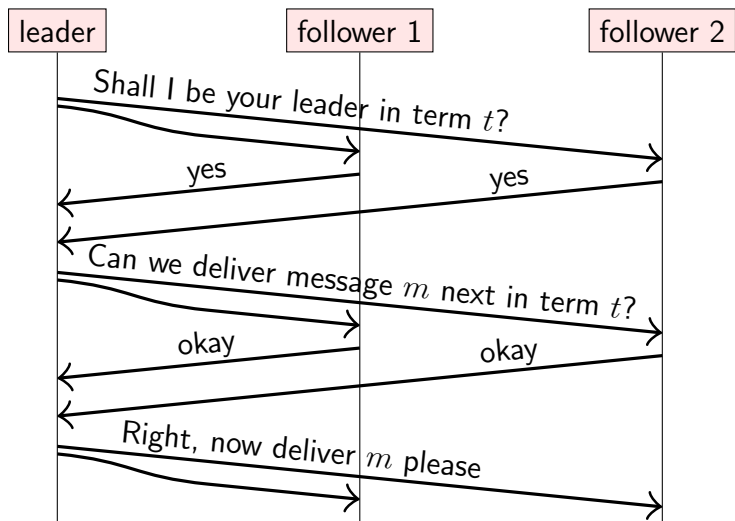
Checking if a leader has been voted out

For every decision (message to deliver), the leader must first get acknowledgements from a quorum.





Checking if a leader has been voted out


For every decision (message to deliver), the leader must first get acknowledgements from a quorum.




Node state transitions in Raft




Node state transitions in Raft



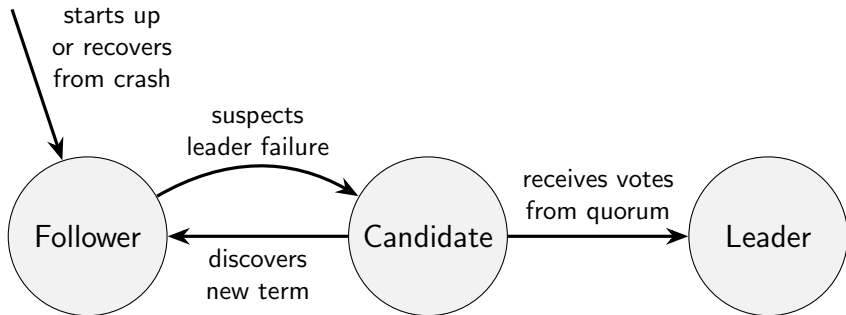
Node state transitions in Raft




Node state transitions in Raft




Node state transitions in Raft



Node state transitions in Raft



Node state transitions in Raft



Raft (1/9): initialisation

on initialisation **do**

currentTerm := 0; *votedFor* := null

log := $\langle \rangle$; *commitLength* := 0

currentRole := follower; *currentLeader* := null

votesReceived := $\{ \}$; *sentLength* := $\langle \rangle$; *ackedLength* := $\langle \rangle$

end on

on recovery from crash **do**

currentRole := follower; *currentLeader* := null

votesReceived := $\{ \}$; *sentLength* := $\langle \rangle$; *ackedLength* := $\langle \rangle$

end on

on node *nodeId* suspects leader has failed, or on election timeout **do**

currentTerm := *currentTerm* + 1; *currentRole* := candidate

votedFor := *nodeId*; *votesReceived* := $\{ nodeId \}$; *lastTerm* := 0

if *log.length* > 0 **then** *lastTerm* := *log*[*log.length* - 1].term; **end if**


msg := (VoteRequest, *nodeId*, *currentTerm*, *log.length*, *lastTerm*)

for each *node* \in *nodes*: **send** *msg* to *node*

start election timer

end on

Raft (1/9): initialisation



on initialisation **do**

$currentTerm := 0$; $votedFor := null$

$log := \langle \rangle$; $commitLength := 0$

$currentRole := follower$; $currentLeader := null$

$votesReceived := \{\}$; $sentLength := \langle \rangle$; $ackedLength := \langle \rangle$

end on

on recovery from crash **do**

$currentRole := follower$; $currentLeader := null$

$votesReceived := \{\}$; $sentLength := \langle \rangle$; $ackedLength := \langle \rangle$

end on

on node $nodeId$ suspects leader has failed, or on election timeout **do**

$currentTerm := currentTerm + 1$; $currentRole := candidate$

$votedFor := nodeId$; $votesReceived := \{nodeId\}$; $lastTerm := 0$

if $log.length > 0$ **then** $lastTerm := log[log.length - 1].term$; **end if**

$msg := (VoteRequest, nodeId, currentTerm, log.length, lastTerm)$

for each $node \in nodes$: **send** msg to $node$

start election timer

end on

Raft (2/9): voting on a new leader

on receiving (VoteRequest, cId , $cTerm$, $cLogLength$, $cLogTerm$)
 at node $nodeId$ **do**

$myLogTerm := log[log.length - 1].term$

$logOk := (cLogTerm > myLogTerm) \vee$
 $(cLogTerm = myLogTerm \wedge cLogLength \geq log.length)$

$termOk := (cTerm > currentTerm) \vee$
 $(cTerm = currentTerm \wedge votedFor \in \{cId, null\})$

if $logOk \wedge termOk$ **then**

$currentTerm := cTerm$

$currentRole := follower$

$votedFor := cId$

send (VoteResponse, $nodeId$, $currentTerm$, true) to node cId

else

send (VoteResponse, $nodeId$, $currentTerm$, false) to node cId

end if

end on

Raft (2/9): voting on a new leader

c for candidate

on receiving (VoteRequest, *cId*, *cTerm*, *cLogLength*, *cLogTerm*)
at node *nodeId* **do**

$myLogTerm := log[log.length - 1].term$

$logOk := (cLogTerm > myLogTerm) \vee$
 $(cLogTerm = myLogTerm \wedge cLogLength \geq log.length)$

$termOk := (cTerm > currentTerm) \vee$
 $(cTerm = currentTerm \wedge votedFor \in \{cId, null\})$

if $logOk \wedge termOk$ **then**

$currentTerm := cTerm$

$currentRole := follower$

$votedFor := cId$

send (VoteResponse, *nodeId*, *currentTerm*, true) to node *cId*

else

send (VoteResponse, *nodeId*, *currentTerm*, false) to node *cId*

end if

end on

Raft (3/9): collecting votes

```
on receiving (VoteResponse, voterId, term, granted) at nodeId do
  if currentRole = candidate  $\wedge$  term = currentTerm  $\wedge$  granted then
    votesReceived := votesReceived  $\cup$  {voterId}
    if |votesReceived|  $\geq \lceil (|nodes| + 1)/2 \rceil$  then
      currentRole := leader; currentLeader := nodeId
      cancel election timer
      for each follower  $\in nodes \setminus \{nodeId\}$  do
        sentLength[follower] := log.length
        ackedLength[follower] := 0
        REPLICATELOG(nodeId, follower)
      end for
    end if
  else if term > currentTerm then
    currentTerm := term
    currentRole := follower
    votedFor := null
    cancel election timer
  end if
end on
```

Raft (4/9): broadcasting messages

```
on request to broadcast msg at node nodeId do
  if currentRole = leader then
    append the record (msg : msg, term : currentTerm) to log
    ackedLength[nodeId] := log.length
    for each follower ∈ nodes \ {nodeId} do
      REPLICATELOG(nodeId, follower)
    end for
  else
    forward the request to currentLeader via a FIFO link
  end if
end on

periodically at node nodeId do
  if currentRole = leader then
    for each follower ∈ nodes \ {nodeId} do
      REPLICATELOG(nodeId, follower)
    end for
  end if
end do
```

Raft (5/9): replicating from leader to followers

Called on the leader whenever there is a new message in the log, and also periodically. If there are no new messages, *entries* is the empty list. LogRequest messages with *entries* = $\langle \rangle$ serve as heartbeats, letting followers know that the leader is still alive.

```
function REPLICATELOG(leaderId, followerId)  
  i := sentLength[followerId]  
  entries :=  $\langle \log[i], \log[i + 1], \dots, \log[\log.length - 1] \rangle$   
  prevLogTerm := 0  
  if i > 0 then  
    prevLogTerm :=  $\log[i - 1].term$   
  end if  
  send (LogRequest, leaderId, currentTerm, i, prevLogTerm,  
    commitLength, entries) to followerId  
end function
```

Raft (6/9): followers receiving messages

```
on receiving (LogRequest, leaderId, term, logLength, logTerm,  
             leaderCommit, entries) at node nodeId do  
  if term > currentTerm then  
    currentTerm := term; votedFor := null  
  end if  
  logOk := (log.length ≥ logLength)  
  if logOk ∧ (logLength > 0) then  
    logOk := (logTerm = log[logLength - 1].term)  
  end if  
  
  if term = currentTerm ∧ logOk then  
    currentRole := follower; currentLeader := leaderId  
    APPENDENTRIES(logLength, leaderCommit, entries)  
    ack := logLength + entries.length  
    send (LogResponse, nodeId, currentTerm, ack, true) to leaderId  
  else  
    send (LogResponse, nodeId, currentTerm, 0, false) to leaderId  
  end if  
end on
```


Raft (7/9): updating followers' logs

```
function APPENDENTRIES(logLength, leaderCommit, entries)  
  if entries.length > 0  $\wedge$  log.length > logLength then  
    if log[logLength].term  $\neq$  entries[0].term then  
      log :=  $\langle$ log[0], log[1], ..., log[logLength - 1] $\rangle$   
    end if  
  end if  
  
  if logLength + entries.length > log.length then  
    for i := log.length - logLength to entries.length - 1 do  
      append entries[i] to log  
    end for  
  end if  
  
  if leaderCommit > commitLength then  
    for i := commitLength to leaderCommit - 1 do  
      deliver log[i].msg to the application  
    end for  
    commitLength := leaderCommit  
  end if  
end function
```

Raft (8/9): leader receiving log acknowledgements

```
on receiving (LogResponse, follower, term, ack, success) at nodeId do  
  if term = currentTerm  $\wedge$  currentRole = leader then  
    if success = true  $\wedge$  ack  $\geq$  ackedLength[follower] then  
      sentLength[follower] := ack  
      ackedLength[follower] := ack  
      COMMITLOGENTRIES()  
    else if sentLength[follower] > 0 then  
      sentLength[follower] := sentLength[follower] - 1  
      REPLICATELOG(nodeId, follower)  
    end if  
  else if term > currentTerm then  
    currentTerm := term  
    currentRole := follower  
    votedFor := null  
  end if  
end on
```

Raft (9/9): leader committing log entries

Any log entries that have been acknowledged by a quorum of nodes are ready to be committed by the leader. When a log entry is committed, its message is delivered to the application.

define $\text{acks}(\text{length}) = |\{n \in \text{nodes} \mid \text{ackedLength}[n] \geq \text{length}\}|$

function COMMITLOGENTRIES

$\text{minAcks} := \lceil (|\text{nodes}| + 1)/2 \rceil$

$\text{ready} := \{len \in \{1, \dots, \text{log.length}\} \mid \text{acks}(len) \geq \text{minAcks}\}$

if $\text{ready} \neq \{\}$ $\wedge \max(\text{ready}) > \text{commitLength} \wedge$
 $\text{log}[\max(\text{ready}) - 1].\text{term} = \text{currentTerm}$ **then**

for $i := \text{commitLength}$ **to** $\max(\text{ready}) - 1$ **do**

 deliver $\text{log}[i].\text{msg}$ to the application

end for

$\text{commitLength} := \max(\text{ready})$

end if

end function